

Nondeterministic Asynchronous Dataflow in Isabelle/HOL

Rafael Castro G. Silva, Laouen Fernet and Dmitriy Traytel

`razi@di.ku.dk`

Department of Computer Science
University of Copenhagen

14/09/2024

Introduction

Time-Aware Stream Processing

- What is Time-Aware Stream Processing?

Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute (possibly) unbounded sequences of data (streams)

Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute (possibly) unbounded sequences of data (streams)
 - Time-Aware: Data has timestamp metadata; timestamps are bounded by watermarks
 - Timestamp: A partially-ordered value associated with the data (e.g., unix-time, int, pairs of ints, etc...)
 - Watermark: A value of the same type of the timestamp. Represents data-completeness.



Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute (possibly) unbounded sequences of data (streams)
 - Time-Aware: Data has timestamp metadata; timestamps are bounded by watermarks
 - Timestamp: A partially-ordered value associated with the data (e.g., unix-time, int, pairs of ints, etc...)
 - Watermark: A value of the same type of the timestamp. Represents data-completeness.




- Asynchronous Dataflow Programming: Directed graph of interconnected operators that perform event-wise transformations
- E.g.: Apache Flink, Apache Samza, Apache Spark, Google Cloud Dataflow, and Timely Dataflow



Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute (possibly) unbounded sequences of data (streams)
 - Time-Aware: Data has timestamp metadata; timestamps are bounded by watermarks
 - Timestamp: A partially-ordered value associated with the data (e.g., unix-time, int, pairs of ints, etc...)
 - Watermark: A value of the same type of the timestamp. Represents data-completeness.



- Asynchronous Dataflow Programming: Directed graph of interconnected operators that perform event-wise transformations
- E.g.: Apache Flink, Apache Samza, Apache Spark, Google Cloud Dataflow, and Timely Dataflow



- Why?
 - Highly Parallel
 - Low latency (output as soon as possible)
 - Incremental computing (re-uses previous computations)

Formalization in Isabelle/HOL

- Classical higher-order logic (HOL): Simple Typed Lambda Calculus + (Hilbert) axiom of choice + axiom of infinity + rank-1 polymorphism

- Classical higher-order logic (HOL): Simple Typed Lambda Calculus + (Hilbert) axiom of choice + axiom of infinity + rank-1 polymorphism
- Isabelle: A generic proof assistant



- Isabelle/HOL: Isabelle's flavor of HOL

- Codatatypes

- Recursion: always eventually reduces an argument
- Corecursion: always eventually produces something

- Recursion: always eventually reduces an argument
- Corecursion: always eventually produces something
- Corec:

- Inductive predicate
 - Finite number of introduction rule applications

- Inductive predicate
 - Finite number of introduction rule applications
- Coinductive predicate
 - Infinite number of introduction rule applications

Isabelle/HOL: (Co)inductive Predicates

- Inductive predicate
 - Finite number of introduction rule applications
- Coinductive predicate
 - Infinite number of introduction rule applications
- Coinduction principle

Conclusion

Conclusion

- Isabelle/HOL has a good tool set to formalize and reason about stream processing:
 - Codatatypes, coinductive predicates, corecursion with friends, reasoning up to friends (congruence),
 - Coinduction up to congruence principle is automatically derived for codatatypes (but not for coinductive principles)
- Next step: Feedback loop

Questions, comments and suggestions