

Verified Time-Aware Stream Processing

Rafael Castro G. Silva

`razi@di.ku.dk`

Department of Computer Science
University of Copenhagen

04/06/2024

Introduction

Time-Aware Stream Processing

- What is Time-Aware Stream Processing?

Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute unbounded sequences of data (streams)

Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute unbounded sequences of data (streams)
 - Time-Aware: Data has timestamp metadata; timestamps are bounded by watermarks
 - Timestamp: A partially-ordered value associated with the data (e.g., unix-time, int, pairs of ints, etc...)
 - Watermark: A value of the same type of the timestamp. Represents data-completeness.



Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute unbounded sequences of data (streams)
 - Time-Aware: Data has timestamp metadata; timestamps are bounded by watermarks
 - Timestamp: A partially-ordered value associated with the data (e.g., unix-time, int, pairs of ints, etc...)
 - Watermark: A value of the same type of the timestamp. Represents data-completeness.



- Asynchronous Dataflow Programming: Directed graph of interconnected operators that perform event-wise transformations
- E.g.: Apache Flink, Apache Samza, Apache Spark, Google Cloud Dataflow, and Timely Dataflow



Time-Aware Stream Processing

- What is Time-Aware Stream Processing?
 - Stream Processing: programs that compute unbounded sequences of data (streams)
 - Time-Aware: Data has timestamp metadata; timestamps are bounded by watermarks
 - Timestamp: A partially-ordered value associated with the data (e.g., unix-time, int, pairs of ints, etc...)
 - Watermark: A value of the same type of the timestamp. Represents data-completeness.



- Asynchronous Dataflow Programming: Directed graph of interconnected operators that perform event-wise transformations
- E.g.: Apache Flink, Apache Samza, Apache Spark, Google Cloud Dataflow, and Timely Dataflow



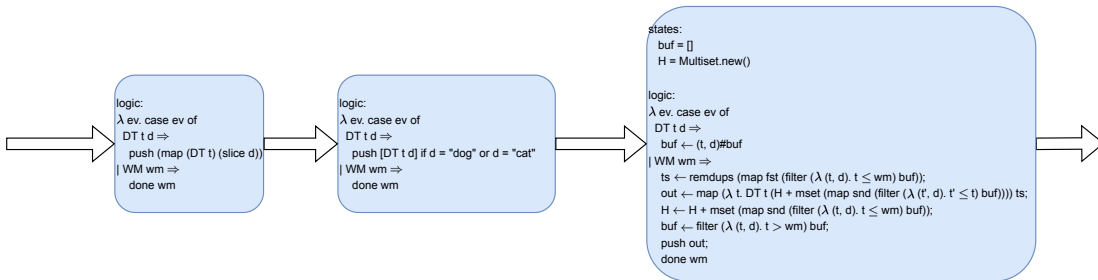
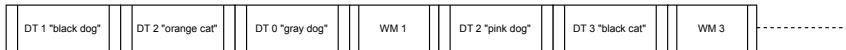
- Why?
 - Highly Parallel
 - Low latency (output as soon as possible)
 - Incremental computing

Time-Aware Stream Processing Example

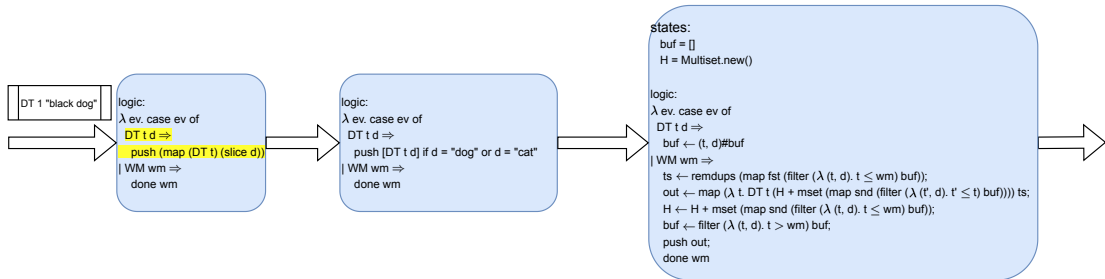
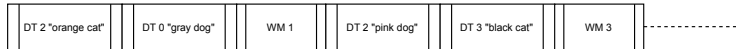
Example:

Incrementally count the occurrences of the words “dog” and “cat”

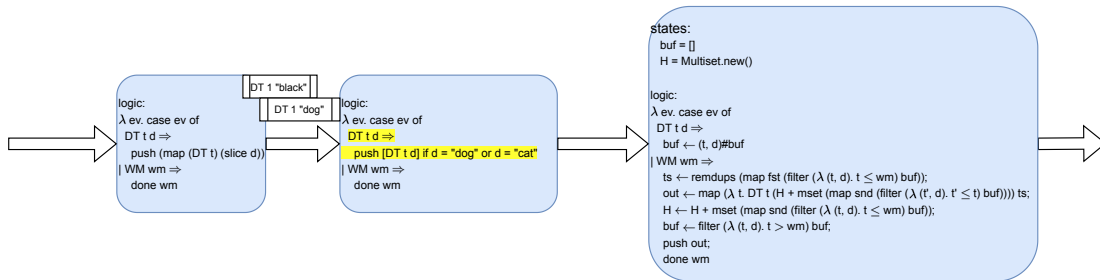
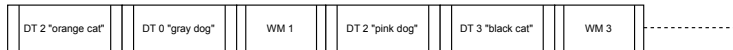
Time-Aware Stream Processing Example



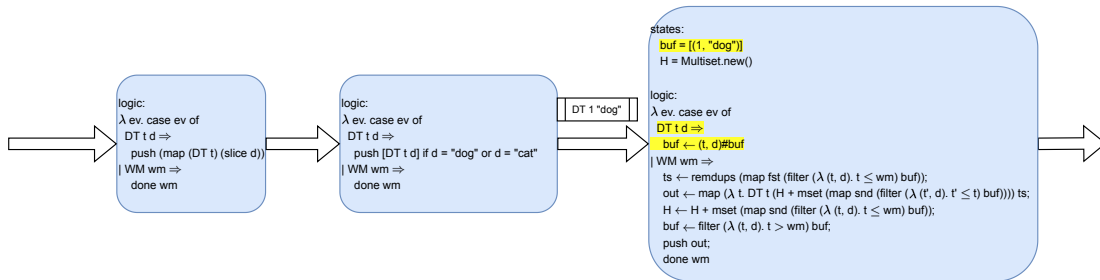
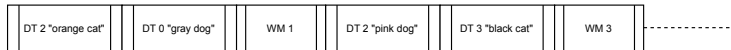
Time-Aware Stream Processing Example



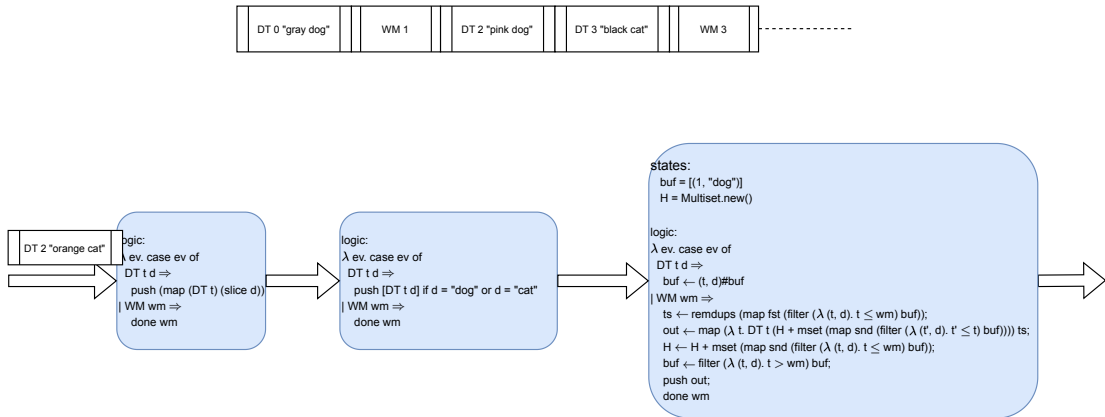
Time-Aware Stream Processing Example



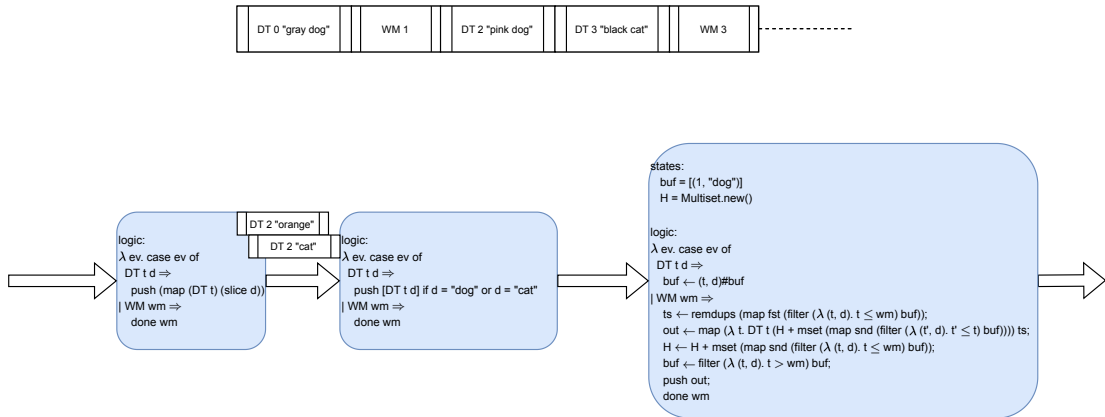
Time-Aware Stream Processing Example



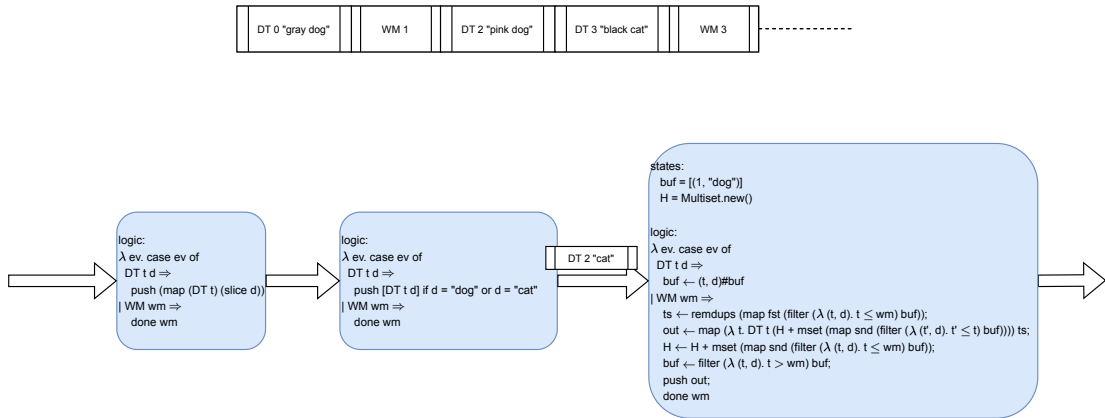
Time-Aware Stream Processing Example



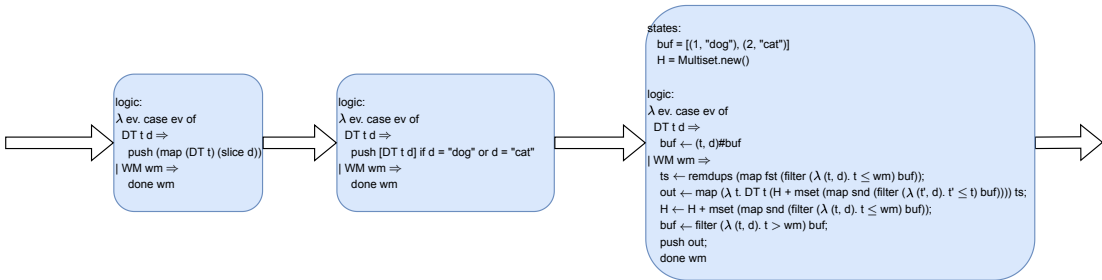
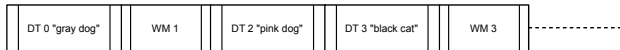
Time-Aware Stream Processing Example



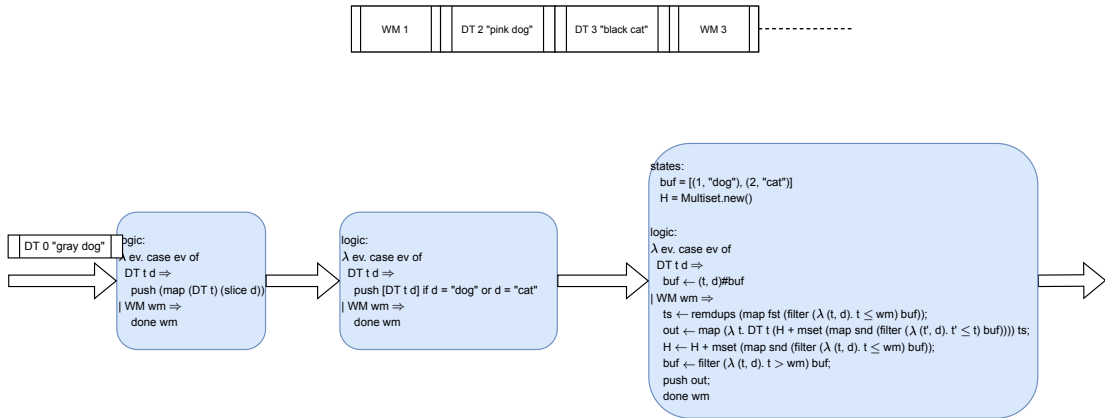
Time-Aware Stream Processing Example



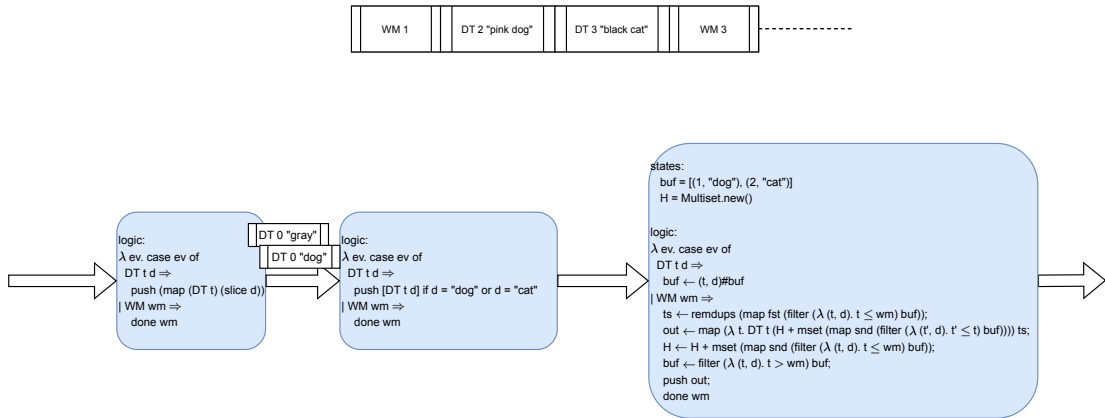
Time-Aware Stream Processing Example



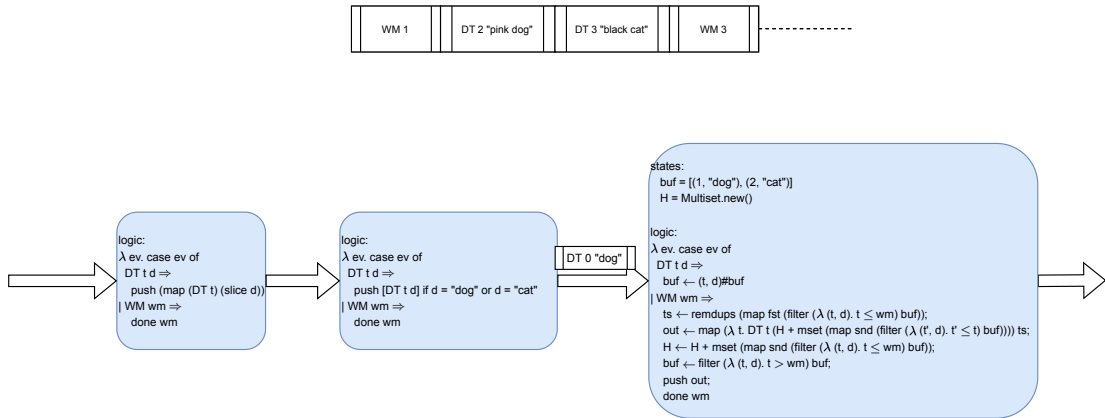
Time-Aware Stream Processing Example



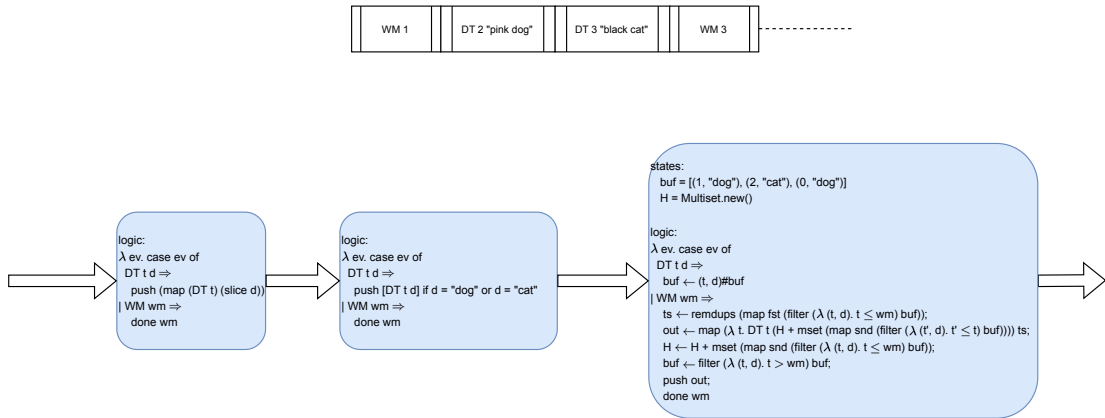
Time-Aware Stream Processing Example



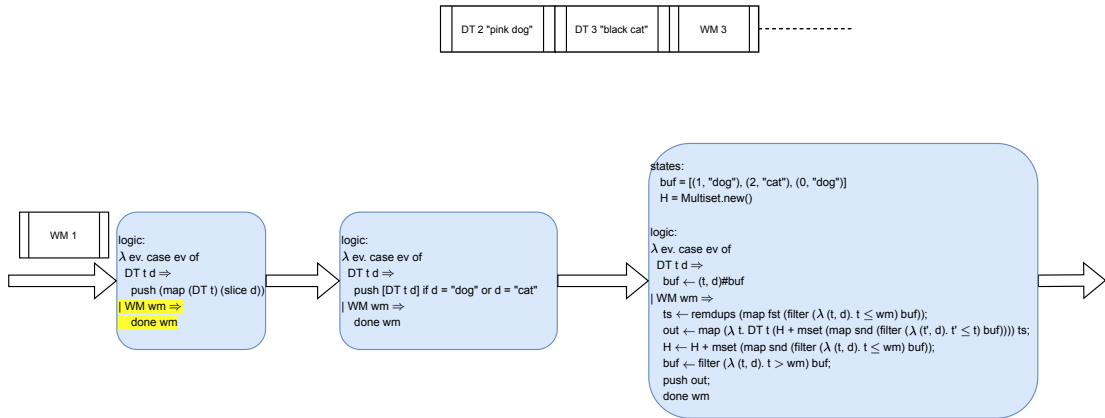
Time-Aware Stream Processing Example



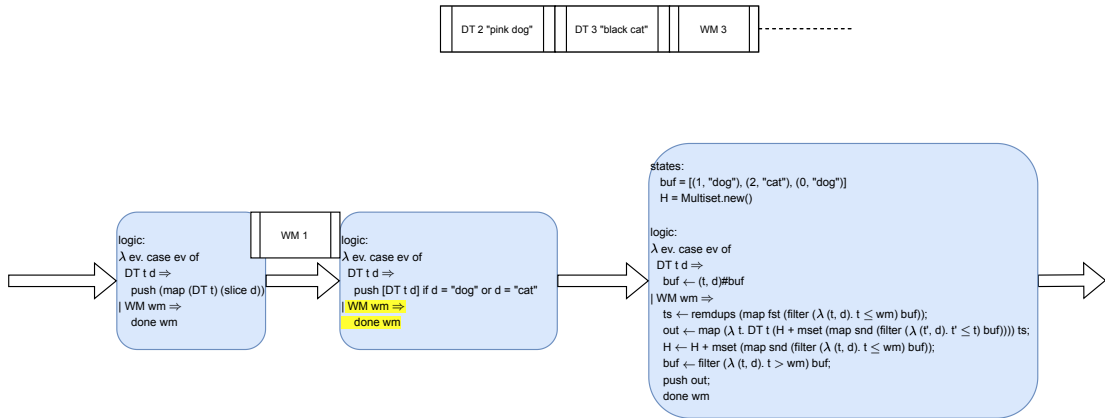
Time-Aware Stream Processing Example



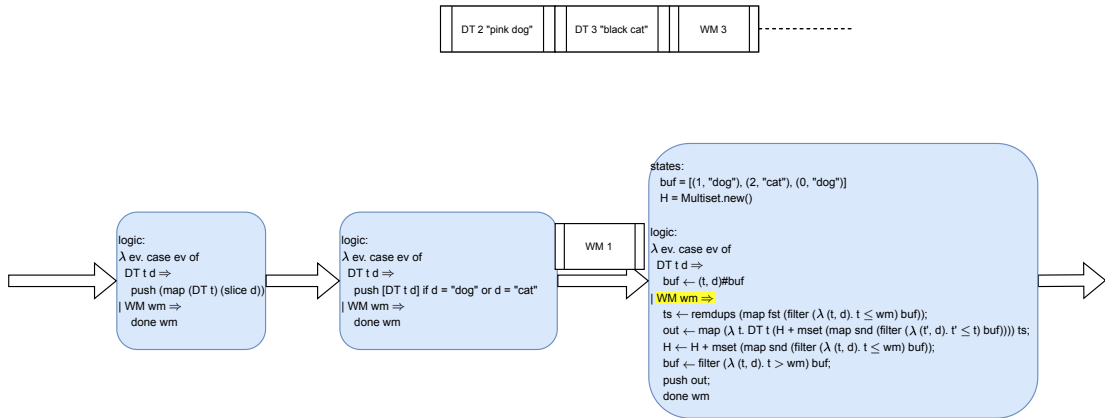
Time-Aware Stream Processing Example



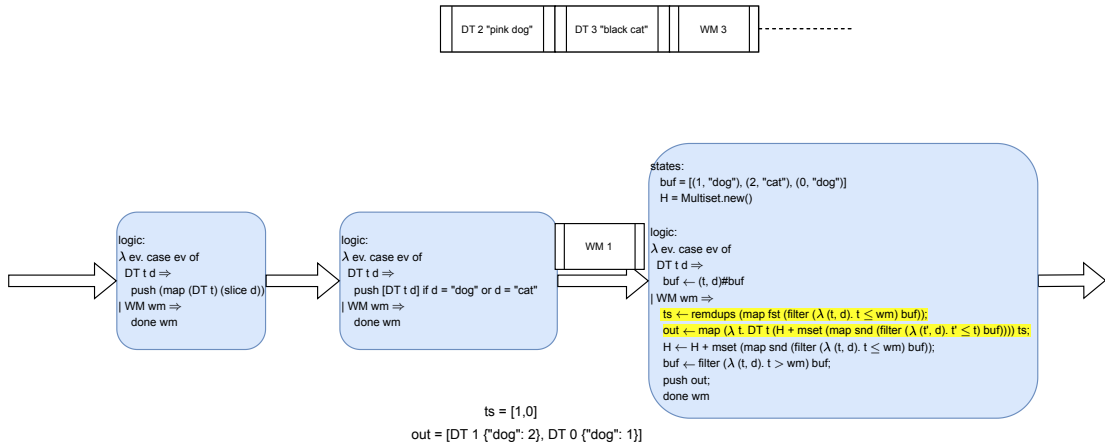
Time-Aware Stream Processing Example



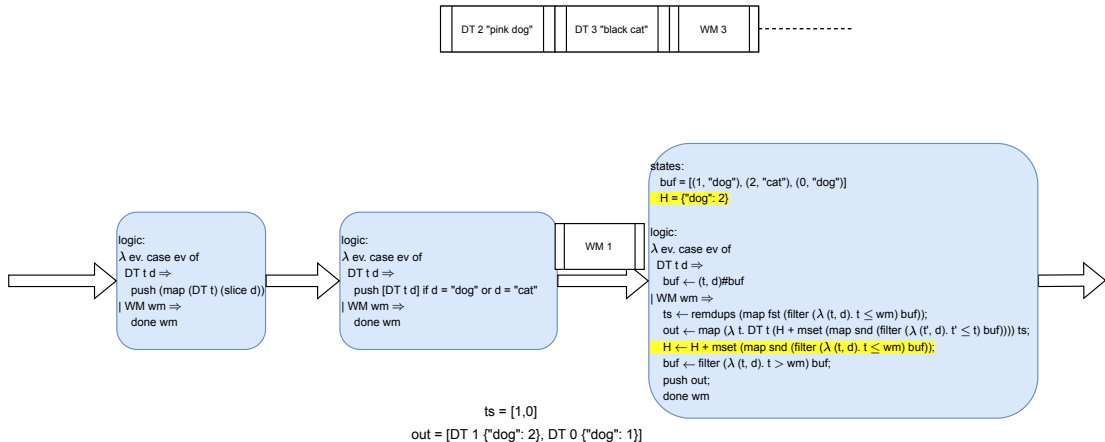
Time-Aware Stream Processing Example



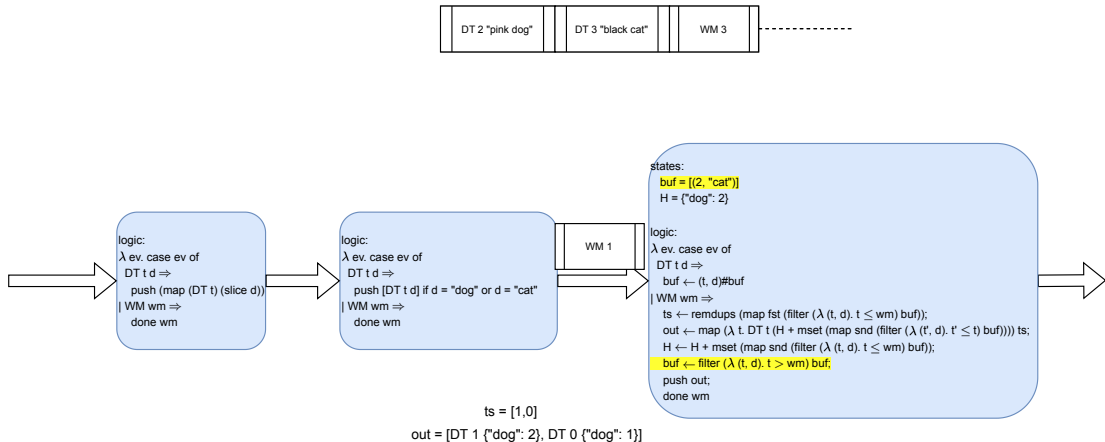
Time-Aware Stream Processing Example



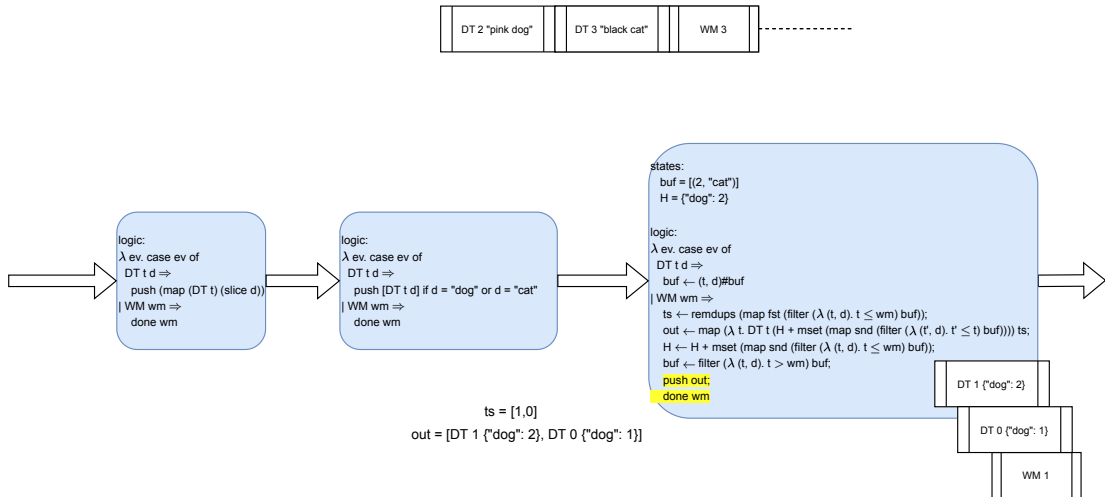
Time-Aware Stream Processing Example



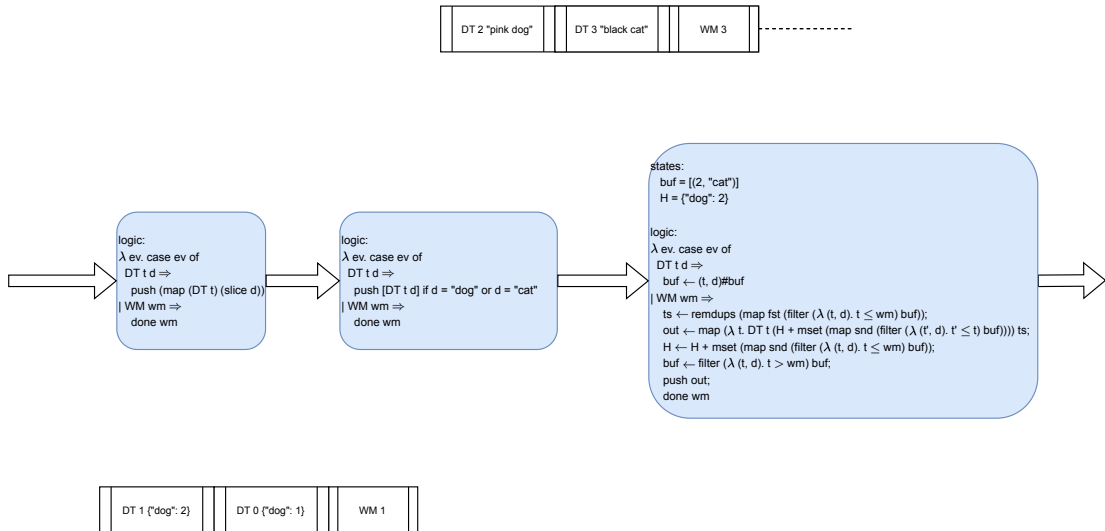
Time-Aware Stream Processing Example



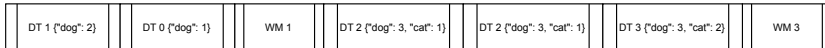
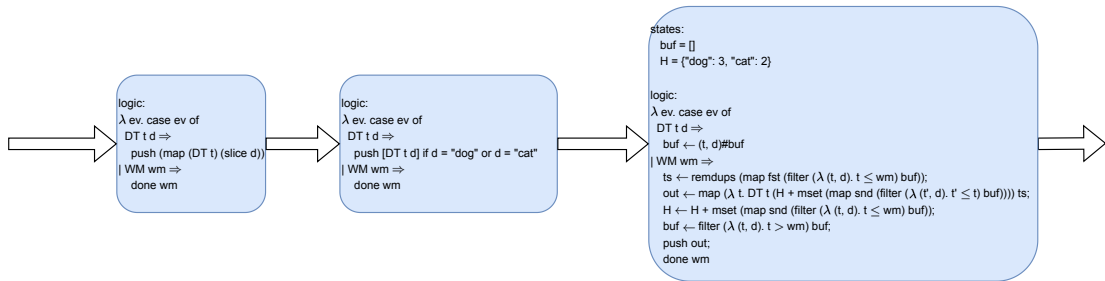
Time-Aware Stream Processing Example



Time-Aware Stream Processing Example



Time-Aware Stream Processing Example



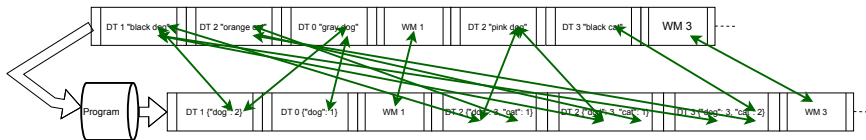
- How do we know if our Dataflow program is what we want?

Properties

- How do we know if our Dataflow program is what we want?
- We need a correctness specification

Properties

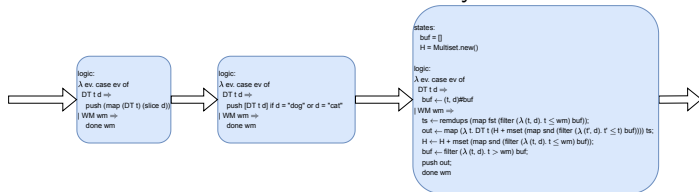
- How do we know if our Dataflow program is what we want?
- We need a correctness specification
- **Intuition of the specification:**
 - Soundness: for every output $DT\ t\ H$, the “dog” count in H is the count of events with timestamp (\leq) t which contains the string “dog”; similarly for “cat”. The count for any other word is always 0.
 - Completeness: The other way around.



How to prove it

- lets break down the problem!:

- The correctness of the entire Dataflow emerges from the correctness of each part (operator)
 - Operator 1: Slicer
 - Operator 2: Filter
 - Operator 3: Incremental histogram
 - Assumptions about the incoming stream:
 - Monotone: after WM w_m no $DT\ t\ d$ such that $t \leq w_m$.
 - Productive: after $DT\ t\ d$ eventually WM w_m such that $t \leq w_m$



- The original incoming stream must respect monotonicity and productivity



- Each operator must preserve monotonicity and productivity!

Writing it down in a proof assistant!

- Classical higher-order logic (HOL): Simple Typed Lambda Calculus + (Hilbert) axiom of choice + axiom of infinity + rank-1 polymorphism

- Classical higher-order logic (HOL): Simple Typed Lambda Calculus + (Hilbert) axiom of choice + axiom of infinity + rank-1 polymorphism
- Isabelle: A generic proof assistant



- Isabelle/HOL: Isabelle's flavor of HOL

- Datatypes and Codatatypes

```
codatatype (lset: 'a) llist = lnull: LNil | LCons (lhd: 'a) (ltl: 'a llist)  
for map: lmap where ltl LNil = LNil
```

- Examples:

- LNil
- LCons 1 (LCons 2 (LCons 3 LNil))
- LCons 0 (LCons 0 (LCons 0 (...)))

- Datatypes and Codatatypes

```
codatatype (lset: 'a) llist = lnull: LNil | LCons (lhd: 'a) (ltl: 'a llist)  
for map: lmap where ltl LNil = LNil
```

- Examples:

- LNil
- LCons 1 (LCons 2 (LCons 3 LNil))
- LCons 0 (LCons 0 (LCons 0 (...)))

- Coinductive principle for lazy list equality:

- Recursion

```
fun lshift :: 'a list  $\Rightarrow$  'a llist  $\Rightarrow$  'a llist (infixr @@ 65) where  
  lshift [] lxs = lxs  
| lshift (x # xs) lxs = LCons x (lshift xs lxs)
```

- While Combinator

```
definition while_option :: ('a  $\Rightarrow$  bool)  $\Rightarrow$  ('a  $\Rightarrow$  'a)  $\Rightarrow$  'a  $\Rightarrow$  'a option where  
  while_option b c s = ...
```

- While rule for invariant reasoning (Hoare-style):
 - There is something that holds before a step; that thing still holds after the step

- Corecursion is like recursion, but instead of always eventually reducing an argument it always eventually produces something

- Corecursion is like recursion, but instead of always eventually reducing an argument it always eventually produces something
- Corec:

```
corec lapp :: 'a llist  $\Rightarrow$  'a llist  $\Rightarrow$  'a llist where  
lapp xs lys = case xs of LNil  $\Rightarrow$  lys | LCons x xs'  $\Rightarrow$  LCons x (lapp xs' lys)
```

Isabelle/HOL: (Co)inductive Predicates

- Inductive predicate
 - Finite number of introduction rule applications

```
inductive in_llist :: 'a  $\Rightarrow$  'a llist  $\Rightarrow$  bool where  
  In_llist: in_llist x (LCons x xs)  
| Next_llist: in_llist x xs  $\Rightarrow$  in_llist x (LCons y xs)  
  
in_llist 2 (LCons 1 (LCons (2 (...))))
```

Isabelle/HOL: (Co)inductive Predicates

- Inductive predicate
 - Finite number of introduction rule applications

```
inductive in_llist :: 'a  $\Rightarrow$  'a llist  $\Rightarrow$  bool where  
  In_llist: in_llist x (LCons x lxs)  
| Next_llist: in_llist x lxs  $\Rightarrow$  in_llist x (LCons y lxs)  
  
in_llist 2 (LCons 1 (LCons (2 (...))))
```

- Coinductive predicate
 - Infinite number of introduction rule applications

```
coinductive lprefix :: 'a llist  $\Rightarrow$  'a llist  $\Rightarrow$  bool where  
  LNil_lprefix: lprefix LNil lxs  
| LCons_lprefix: lprefix lxs lxs  $\Rightarrow$  lprefix (LCons x lxs) (LCons x lxs)  
  
lprefix (LCons 1 (LCons (2 (...)))) (LCons 1 (LCons (2 (...))))
```

Isabelle/HOL: (Co)inductive Predicates

- Inductive predicate
 - Finite number of introduction rule applications

```
inductive in_llist :: 'a  $\Rightarrow$  'a llist  $\Rightarrow$  bool where  
  In_llist: in_llist x (LCons x lxs)  
| Next_llist: in_llist x lxs  $\Rightarrow$  in_llist x (LCons y lxs)  
  
in_llist 2 (LCons 1 (LCons (2 (...))))
```

- Coinductive predicate
 - Infinite number of introduction rule applications

```
coinductive lprefix :: 'a llist  $\Rightarrow$  'a llist  $\Rightarrow$  bool where  
  LNil_lprefix: lprefix LNil lxs  
| LCons_lprefix: lprefix lxs lxs  $\Rightarrow$  lprefix (LCons x lxs) (LCons x lxs)  
  
lprefix (LCons 1 (LCons (2 (...)))) (LCons 1 (LCons (2 (...))))
```

- Coinduction principle

Lazy Lists Processors

Operator formalization

- Operator as a codatatype

- Taking $'i$ as the input type, and $'o$ as the output type:

$\text{codatatype } ('o, 'i) \text{ op} = \text{Logic } (\text{apply}: ('i \Rightarrow ('o, 'i) \text{ op} \times 'o \text{ list}))$

Operator formalization

- Operator as a codatatype
 - Taking `'i` as the input type, and `'o` as the output type:
`codatatype ('o, 'i) op = Logic (apply: ('i \Rightarrow ('o, 'i) op \times 'o list))`
 - Infinite trees: applying the selector `apply` “walks” a branch of the tree

- Produce function: applies the logic (co)recursively throughout a lazy list

definition $\text{produce}_1 \text{ op } lxs = \text{while_option} \dots$

corec produce **where**

produce $\text{op } lxs = (\text{case } \text{produce}_1 \text{ op } lxs \text{ of}$

None \Rightarrow LNil

| Some $(\text{op}', x, xs, lxs') \Rightarrow \text{LCons } x (xs @@ \text{produce } \text{op}' lxs')$)

- Produce function: applies the logic (co)recursively throughout a lazy list

definition $\text{produce}_1 \text{ op } lxs = \text{while_option} \dots$

corec produce **where**

$\text{produce } op \text{ } lxs = (\text{case } \text{produce}_1 \text{ op } lxs \text{ of}$

None \Rightarrow LNil

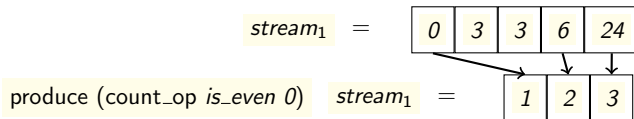
| Some $(op', x, xs, lxs') \Rightarrow \text{LCons } x \text{ } (xs \text{ @@ } \text{produce } op' \text{ } lxs'))$

- produce_1 has an induction principle based on the while invariant rule

Operators: Count

- Example:

```
corec count_op where count_op P n =  
  Logic (λe. if P e then (count_op P (n + 1), [n+1]) else (count_op P n, []))
```



Sequential Composition operator

- Sequential composition: take the output of the first operator and give it as input to the second operator.

- Correctness:

$$\text{produce } (\text{comp_op } op_1 \ op_2) \ xs = \text{produce } op_2 \ (\text{produce } op_1 \ xs)$$

- Proof: coinduction principle for lazy list equality and produce_1 induction principle

Time-Aware Operators

- Time-Aware lazy lists

```
datatype ('t::order, 'd) event = DT (tmp: 't) (data: 'd) | WM (wmk: 't)
```

- Time-Aware lazy lists

```
datatype ('t::order, 'd) event = DT (tmp: 't) (data: 'd) | WM (wmk: 't)
```

- Generalization to partial orders
 - Cycles
 - Operators with multiple inputs
- Productive and monotone streams: Coinductive predicates over lazy lists of events.

Proving histogram correct: building Blocks

- Histogram operator: batching and incremental computing
- Building Blocks: reusable operators
 - Batching: `batch_op`
 - Incremental computing: `incr_op`
 - Soundness, completeness, preservation of monotonicity and productivity

- `batch_op` and `incr_op` can be composed

definition $\text{incr_batch_op } buf1 \text{ } buf2 = \text{comp_op } (\text{batch_op } buf1) (\text{incr_op } buf2)$

- Soundness, completeness, and monotone and productive preservation

Histogram Operator

```
corec map_op where map_op f = Logic ( $\lambda$  ev. case ev of  
WM wm  $\Rightarrow$  (map_op f, [WM wm]) | DT t d  $\Rightarrow$  (map_op f, [DT t (f t d)]))
```

```
definition incr_hist_op buf1 buf2 =  
comp_op (incr_batch_op buf1 buf2) (map_op incr_coll)
```

- Soundness, completeness, and monotone and productive preservation

Other shapes

Join Operator

- Relation Join
- Use the `sum` type in the timestamps to represent two stream as one
- Partial order for the `sum`: lefts and rights are incomparable
- Defined using `incr_batch_op`
- Soundness, completeness, and monotone and productive preservation

Next Steps

Next Steps

- Feedback loop

Questions, comments and suggestions